

# Applied Deep Learning

DAT310

**Lecture:**  
**Deep Learning Best**  
**Practices**



UNIVERSITY OF  
CALGARY

**Start**  
**something.**

# Agenda



UNIVERSITY OF  
CALGARY

**Start  
something.**

- General ML Pipeline
- Rules of Thumb
- Debugging Networks
- Improving Results
- Faster Networks



# General ML Pipeline



**UNIVERSITY OF CALGARY**  
CONTINUING EDUCATION

**Start  
something.**

# General ML Pipeline



UNIVERSITY OF  
CALGARY

**Start  
something.**

1. Get a dataset
2. Build a dataset and dataloader
3. Choose some models for the task
4. Train your models
  - a. Need optimizer and loss function
  - b. Compute loss w/ output of network, compute gradients using backwards
  - c. Apply gradients with optimizer.step()
5. Evaluate your models



# Rules of Thumb



**UNIVERSITY OF CALGARY**  
CONTINUING EDUCATION

**Start  
something.**

# Rules of Thumb




UNIVERSITY OF  
CALGARY

Start  
something.

## Optimizer

- Pick Adam as your optimizer
  - Combo of Momentum and RMS Prop
  - Automatically figures out a learning rate for you!



# Rules of Thumb




UNIVERSITY OF  
CALGARY

Start  
something.

## Non Linearities

- Deeper VS Wider -> Deeper is better
  - Why? More non linearities
- How to make deeper?
  - Skip Connections



# Rules of Thumb




UNIVERSITY OF  
CALGARY



Start  
something.

## Non Linearities

- Pick Relu
  - No exploding or vanishing gradients like tanh or sigmoid



## Some Common Activation Functions & Their Derivatives





## CV Architecture Choice

- For Computer Vision
  - CNNs -> Bias towards texture
  - Transformers -> Bias towards shape -> Model local visual structures


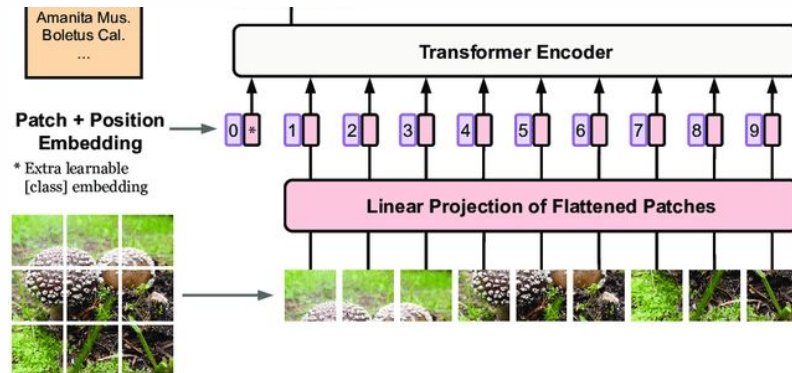



Figure 1: Classification of a standard ResNet-50 of (a) a texture image (elephant skin: only texture cues); (b) a normal image of a cat (with both shape and texture cues), and (c) an image with a texture-shape cue conflict, generated by style transfer between the first two images.



## CV Architecture Choice

- Solution?
  - CNN Features into Transformer



# Rules of Thumb

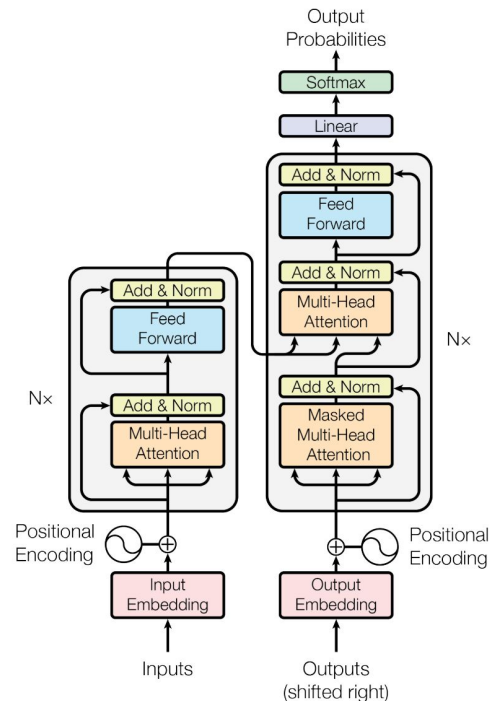


UNIVERSITY OF  
CALGARY

Start  
something.

## NLP Architecture Choice

- No one uses RNNs anymore
- LSTM cells are better for long running sequences
- GRUs better for short
- Transformers better overall
  - Pick a transformer



# Rules of Thumb




UNIVERSITY OF  
CALGARY

Start  
something.

## Reinforcement Learning

- When do you TYPICALLY use reinforcement learning?
  - No loss function with a gradient available
- You have an environment that produces a state
  - I.e a video game



# Rules of Thumb



UNIVERSITY OF  
CALGARY

Start  
something.

## Batch Size

- Bigger = Faster training, worse generalization
- Smaller = Slower training, better generalization



Testing loss and accuracy when the model is trained using different batch sizes.

- Purple = Batch Size 1024
- Blue = 256
- Orange = 64

# Rules of Thumb




UNIVERSITY OF  
CALGARY

Start  
something.

## Learning Rate

- Plot your loss per LR



# Rules of Thumb




UNIVERSITY OF  
CALGARY

Start  
something.

## Learning Rate

- Plot your loss per LR





# Rules of Thumb



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Transfer Learning


- Always use pretrained models!
  - ImageNet or COCO for CV
  - Language Models for NLP





## BatchNorm

- Always use BatchNorm (or some variation)!
  - Trains 6X faster
  - Less sensitive to LR
  - Helps w/ vanishing and exploding gradients



# Rules of Thumb



UNIVERSITY OF  
CALGARY

Start  
something.

## Model Size VS Dataset Size

- If you are positive your implementation of the network and training is correct, but you still cannot overfit to just a small subset of training data.... Make your model deeper + wider
- As long as your training error is low, you can always decrease generalization error by collecting more training data



# Rules of Thumb



UNIVERSITY OF  
CALGARY

**Start  
something.**



## Output Activation

- Binary Classification: Sigmoid
- Multi Class: Softmax
- Regression: Whatever matches the scale of your outputs



## Imbalanced Data

- Oversample
- Undersample
- Weight the loss
- Use balanced accuracy and a confusion matrix to evaluate with



# Rules of Thumb



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Regularization

- L2 is less forgiving, set  $\lambda$  small i.e 0.0001
- L2 is used by default in pytorch



# Rules of Thumb



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Gradient Clipping

- I loathe gradient clipping
- Prevents exploding grads
- Should fix exploding grads before using this
- Prevents gradients from being over a certain size
  - I find this cripples your learning

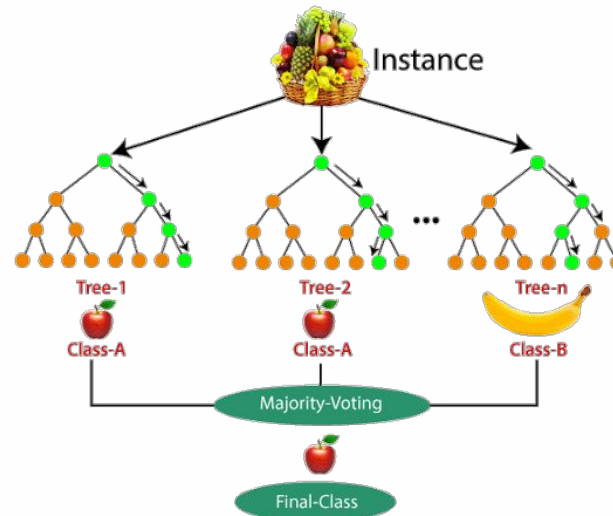


# Rules of Thumb



Don't forget about Classic ML Methods

- Linear Regression, Logistic Regression, Random Forest, Gradient Boosting... all excellent models that are fast to train
- IME they do very well on tabular data, awful on text and images





# Rules of Thumb

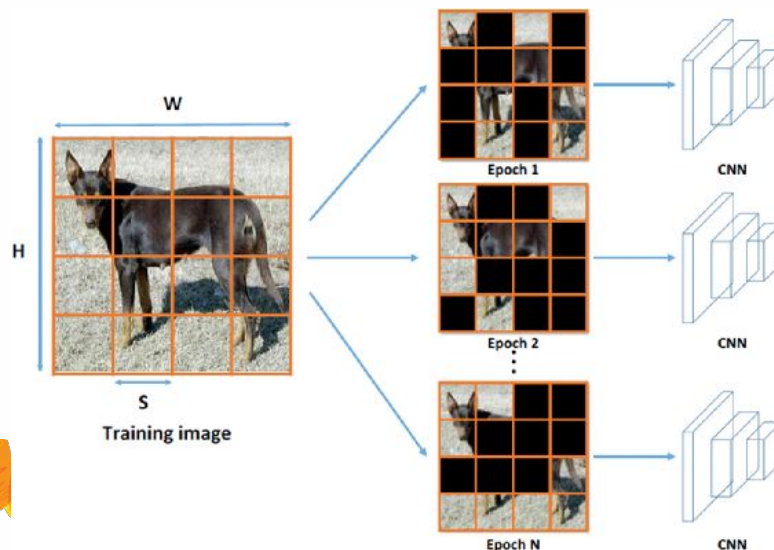


UNIVERSITY OF  
CALGARY

Start  
something.

Look into NON Standard Data Augmentations

- CutMix
- Hide and Seek



	ResNet-50	Mixup [48]	Cutout [3]	CutMix
Image				
Label	Dog 1.0	Dog 0.5 Cat 0.5	Dog 1.0	Dog 0.6 Cat 0.4
ImageNet Cls (%)	76.3 (+0.0)	77.4 (+1.1)	77.1 (+0.8)	<b>78.6</b> (+2.3)
ImageNet Loc (%)	46.3 (+0.0)	45.8 (-0.5)	46.7 (+0.4)	<b>47.3</b> (+1.0)
Pascal VOC Det (mAP)	75.6 (+0.0)	73.9 (-1.7)	75.1 (-0.5)	<b>76.7</b> (+1.1)



# Debugging Networks



**UNIVERSITY OF CALGARY**  
CONTINUING EDUCATION

**Start  
something.**

# Debugging Networks



UNIVERSITY OF  
CALGARY

**Start  
something.**

1. Check dataloader
2. Overfit on a small subset of data
3. Play with LR
4. Are your parameters registered?
5. Check gradients



# Debugging Networks



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Check Dataloader

- For CV insure your images are loaded properly
- Insure you're labels are aligned
- If you apply augmentations to your images, you MIGHT need to apply them to the labels
  - During segmentation for example



# Debugging Networks



UNIVERSITY OF  
CALGARY

Start  
something.

## Overfit on a Small Subset

- Your network should be able to memorize a small subset of data
- This is MUCH faster than using the entire dataset
- Allows you to check that the network works, LR is correct, not too much regularization, etc....  
QUICKLY



# Debugging Networks



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Play With LR

- Covered the plotting a lot already
- Make the plots!



# Debugging Networks



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Check Parameters are Registered

- Print out the parameters of your model
- Make sure all layers are there
- For parameter in model.parameters():  
    print(parameter)



# Debugging Networks




UNIVERSITY OF  
CALGARY

Start  
something.

## Check Gradients

- Can use tensorboardX to viz gradients
- Can check if exploding or vanishing
- Can check if proper size throughout the network



# Improving Performance



**UNIVERSITY OF CALGARY**  
CONTINUING EDUCATION

**Start  
something.**



# Improving Performance



UNIVERSITY OF  
CALGARY

Start  
something.

## Hyperparameter Search

The learning rate is the most important hyperparameter. If bound by time, focus on tuning it.

- The learning rate can be picked by monitoring learning curves that plot the objective function over time.
- The optimal learning rate is typically higher than the learning rate that yields the best performance after the first ~100 iterations, but not so high that it causes instability..



# Improving Performance




UNIVERSITY OF  
CALGARY

Start  
something.

## Hyperparameter Search

Random search typically converges to good hyperparameters faster than grid search.



# Improving Performance



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Check Examples w/ Largest Loss Values

- Save losses on validation or training
- Check the examples with the largest loss values
- Is there a pattern?
- If so, look for papers that address this, or get more data labelled!



# Improving Performance



UNIVERSITY OF  
CALGARY

**Start  
something.**

## Early Stopping

- Compute validation loss + accuracy every epoch
- Save the best epoch, even if it's not the last one
- That's your model!



# Improving Performance

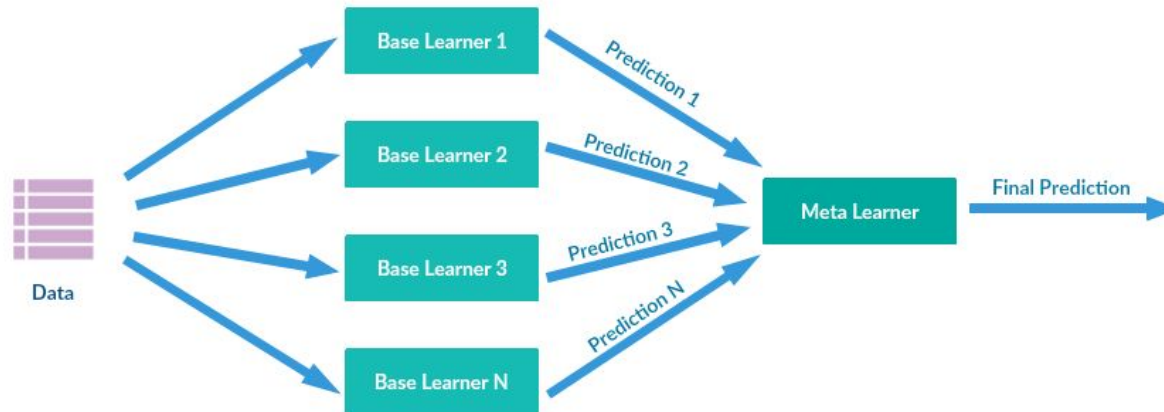


UNIVERSITY OF  
CALGARY

Start  
something.

## Ensembling

- Train multiple models and combine their predictions
  - Usually through averaging the output probabilities
- Great for increasing performance, bad in actual production
  - Too expensive to run!



# Improving Performance



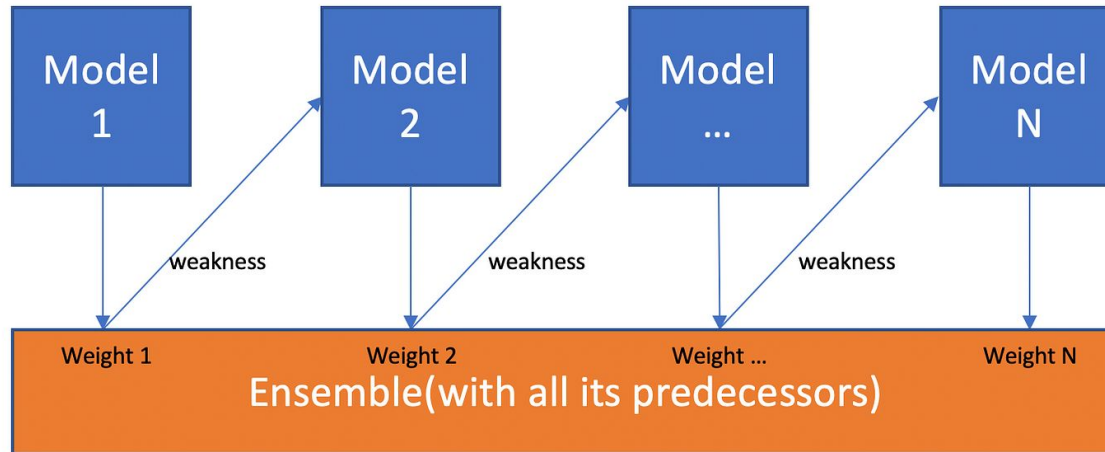
UNIVERSITY OF  
CALGARY

Start  
something.

## Boosting

- Boosting = Training a model on the errors of your previous model
- Then ensembling the two models
- Fancier version of ensembling

Model 1,2,..., N are individual models (e.g. decision tree)



# Improving Performance



UNIVERSITY OF  
CALGARY

Start  
something.

## Increase Resolution

- Higher input resolution almost ALWAYS gives better results
- Can test / validate with a higher input resolution than trained with
  - Form of regularization



# Improving Performance

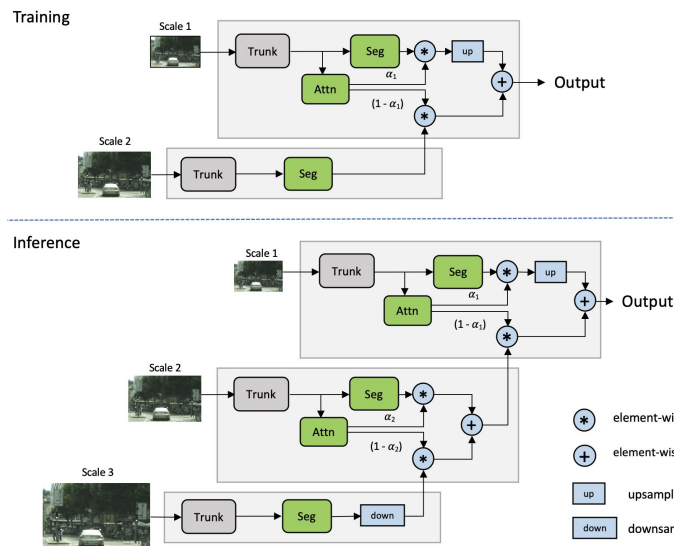


UNIVERSITY OF  
CALGARY

Start  
something.

## Multiscale Predictions

- Train and test using multiple resolutions
- Combine the results of multiple resolutions for your predictions





# Improving Performance




UNIVERSITY OF  
CALGARY

Start  
something.

Use output of all CNN blocks for Predictions

- Concat the output from the end of each block, then make predictions using that



# Improving Performance

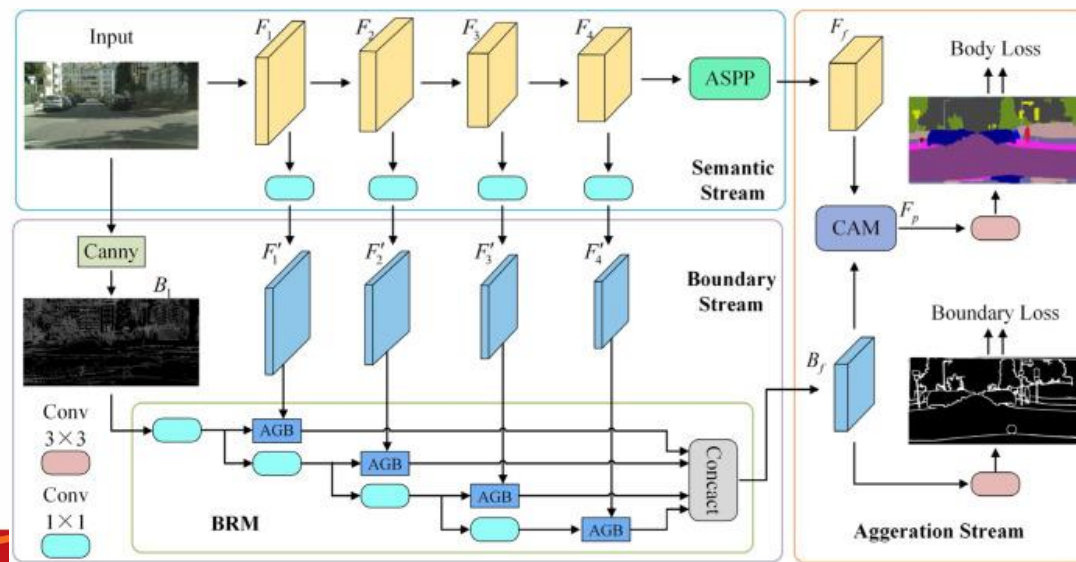


UNIVERSITY OF  
CALGARY

Start  
something.

## Auxiliary Losses

- Use a loss that addresses the weakness of your current model
  - Boundary loss + Segmentation loss for example



# Improving Performance




UNIVERSITY OF  
CALGARY

Start  
something.

Don't just use RGB Input

- Can use RGB + Depth + Edges + a Segmentation etc....
- Sometimes adding this extra channel improves performance



# Improving Performance




UNIVERSITY OF  
CALGARY

Start  
something.

## When Transformers Excel

- Need AT LEAST 10 MILLION training examples
- Not much better than CNNs until over 100 MILLION training examples



# Faster Networks



**UNIVERSITY OF CALGARY**  
CONTINUING EDUCATION

**Start  
something.**

## Pin Memory

- `Dataloader(dataset, pin_memory=True)`
- Pin\_memory transfers to the GPU faster


Setting for the training DataLoader	Time for one training epoch
<code>{'num_workers': 0, 'pin_memory': False}</code>	8.2 s
<code>{'num_workers': 1, 'pin_memory': False}</code>	6.75 s
<code>{'num_workers': 1, 'pin_memory': True}</code>	6.7 s
<code>{'num_workers': 2, 'pin_memory': True}</code>	4.2 s
<code>{'num_workers': 4, 'pin_memory': False}</code>	4.5 s
<code>{'num_workers': 4, 'pin_memory': True}</code>	4.1 s
<code>{'num_workers': 8, 'pin_memory': True}</code>	4.5 s



## Half Precision

- By default Pytorch uses float 32
- Half precision is faster, and uses half as much memory
  - `model.half()`
- Does take a minor performance (accuracy) hit
  - Very minor

## SMALLER AND FASTER



[developer.nvidia.com/tensorrt](https://developer.nvidia.com/tensorrt)

ResNet50 Model, Batch Size = 128, TensorRT 2.1 RC prerelease

19  NVIDIA

## Zero Grads Properly

```
model.zero_grad()
```

```
# or
```

```
optimizer.zero_grad()
```

- executes `memset` for every parameter in the model
- backward pass updates gradients with `"+="` operator (read + write)



```
for param in model.parameters():  
    param.grad = None
```

```
# or (in PyT >= 1.7)
```

```
model.zero_grad(set_to_none=True)
```

- doesn't execute `memset` for every parameter
- memory is zeroed-out by the allocator in a more efficient way
- backward pass updates gradients with `"="` operator (write)



## cuDNN Benchmark

For convolutional neural networks, enable cuDNN autotuner by setting:

```
torch.backends.cudnn.benchmark = True
```

- ▶ cuDNN supports many algorithms to compute convolution
- ▶ autotuner runs a short benchmark and selects algorithm with the best performance

### Example:

[nn.Conv2d](#) with 64 3x3 filters applied to an input with batch size = 32, channels = width = height = 64.

Setting	<code>cudnn.benchmark = False</code> (the default)	<code>cudnn.benchmark = True</code>	Speedup
Forward propagation (FP32) [us]	1430	840	1.70
Forward + backward propagation (FP32) [us]	2870	2260	1.27

# Faster Networks




UNIVERSITY OF  
CALGARY

Start  
something.

## Pytorch 2.0

- Pytorch 2.0 allows you to compile your model
- `Model = torch.compile(model)`
- Provides a massive speed up



Speedups for torch.compile against eager mode on an NVIDIA A100 GPU